

SHARED EVIDENCE CORPUS



13 IAEA core documents



+ 51 same-domain supplements

CONTROLLED EXPANSION



0 5 10 ... 51

$$\mathcal{D}_{m,t} = \mathcal{D}_{core} \cup \mathcal{D}_{sup}^{(m,t)}$$

MANUALLY FIXED QUERIES

Baseline suite

4 bilingual core / phrase queries

P0 extension

6 bilingual technical-detail queries
spikes · zeros · synchrony

No LLM-generated queries
No pooled cosine scale

Shared corpus +
controlled expansion

top-*k* retrieval
k = 10

HISTORICAL

legacy robustness

- nomic index
- 17 indexes
- 100 draws

P0 GRID

primary comparison

- two MiniLMs
- 3 chunk grids
- 10 draws

FIXED WORD

cross-check

- two MiniLMs
- 240 / 24 words
- 10 draws

RECORDED OUTCOMES AND INTERPRETATION BOUNDARY

cosine similarity
ranked within a protocol

IAEA Priority@10
source-policy adherence

first-IAEA rank
authority visibility

strict / document recall
historical self-retrieval only

Do not infer fault accuracy, universal model rankings, or generation quality.